

NeuralNet_Eval copy

June 29, 2023

1 Trabajo Práctico 2: Gradiente Descendente

1.1 Alumnos:

1.1.1 - Luca Mazzarello

1.1.2 - Ignacio Pardo

1.2 Importación de Librerías

```
[ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from tqdm import tqdm
import networkx as nx
import itertools
import seaborn as sns
```

```
[ ]: import warnings

#suppress warnings
warnings.filterwarnings('ignore')
```

```
[ ]: # Display loss with all digits (not rounded)

pd.options.display.float_format = '{:,.20f}'.format
```

1.3 Lectura de Datos

```
[ ]: real_estate: pd.DataFrame = pd.read_excel('data/Real estate valuation data set.
↪xlsx')
real_estate.drop(['No'], axis=1, inplace=True)

short = real_estate.copy()
short.columns = short.columns.map(lambda x: x.split(' ')[0])
display(short.head())
```

	X1	X2	X3 \
0	2,012.9166666999995065155	32.00000000000000000000	84.878820000000000459750
1	2,012.9166666999995065155	19.50000000000000000000	306.5946999999998890416

```

2 2,013.583333330000004934845 13.3000000000000071054 561.98450000000002546585
3 2,013.5000000000000000000000 13.3000000000000071054 561.98450000000002546585
4 2,012.833333330000004934845 5.00000000000000000000 390.5683999999999690772

```

	X4	X5	X6	Y
0	10	24.98298000000000129717	121.54023999999999716692	37.89999999999999857891
1	9	24.980340000000000176556	121.539510000000000703949	42.200000000000000284217
2	5	24.987459999999999867214	121.54390999999999678494	47.29999999999999715783
3	5	24.987459999999999867214	121.54390999999999678494	54.79999999999999715783
4	5	24.97936999999999940769	121.542450000000000231921	43.100000000000000142109

```

[ ]: X_matrix: np.ndarray = np.array(real_estate.drop(['Y house price of unit_
↪area'], axis=1))
y_matrix: np.ndarray = np.array(real_estate['Y house price of unit area']).
↪reshape(-1, 1)

X_train: np.ndarray = X_matrix[0:315]
y_train: np.ndarray = y_matrix[0:315]

X_test: np.ndarray = X_matrix[315:]
y_test: np.ndarray = y_matrix[315:]

X_train.shape, y_train.shape, X_test.shape, y_test.shape

```

```

[ ]: ((315, 6), (315, 1), (99, 6), (99, 1))

```

1.4 Definición de clase NeuralNet

```

[ ]: from typing import List, Tuple, Callable
from matplotlib import pyplot as plt
import numpy as np
from numpy import ndarray
import networkx as nx

class NeuralNet:
    def __init__(self):
        """
        Inicializa los parámetros de la red neuronal.
        2 capas con 5 neuronas cada una.
        """

        self.W1: ndarray = np.random.random((5, 6))
        self.b1: ndarray = np.random.random((5, 1))

        self.W2: ndarray = np.random.random((1, 5))
        self.b2: ndarray = np.random.random((1, 1))

```

```

self.training_loss_acum: List[float] = []
self.testing_loss_acum: List[float] = []

self.activation = self.sigmoid

def sigmoid(self, x: ndarray) -> ndarray:
    """
    Función de activación sigmoide.
    """
    return 1 / (1 + np.exp(-x))

def relu(self, x: ndarray) -> ndarray:
    """
    Función de activación ReLU.
    """
    return np.maximum(0, x)

def leaky_relu(self, x: ndarray) -> ndarray:
    """
    Función de activación Leaky ReLU.
    """
    return np.maximum(0.01 * x, x)

def tanh(self, x: ndarray) -> ndarray:
    """
    Función de activación tanh.
    """
    return np.tanh(x)

def softmax(self, x: ndarray) -> ndarray:
    """
    Función de activación softmax.
    """
    return np.exp(x) / np.sum(np.exp(x))

def forward(
    self,
    x: ndarray,
    dropout: bool = False,
) -> ndarray:
    """
    Calcula la salida de la red neuronal.
    """

    self.z1 = self.W1 @ x.T + self.b1.repeat(x.shape[0], axis=1)
    self.a1 = self.activation(self.z1)

```

```

# if dropout:
#     self.a1 = self.dropout(self.a1, 0.5)

self.z2 = self.W2 @ self.a1 + self.b2.repeat(self.a1.shape[1], axis=1)
self.a2 = self.z2
return self.a2.reshape(-1, 1)

def numerical_gradient(
    self, x: ndarray, y: ndarray, eps: float, dropout: bool = False,
    dropout_prob: float = 0.5
) -> Tuple[ndarray, ndarray, ndarray, ndarray]:
    """
    Una estrategia para calcular estas derivadas parciales,
    consiste en calcular el promedio de los cocientes incrementales
    a derecha e izquierda.
    Para obtener la siguiente aproximación, para
    cada parámetro de la red calculamos:

    $$
    \frac{\partial L}{\partial p} \sim \frac{L(\theta_t, p + \epsilon) -
    \hookrightarrow L(\theta_t, p - \epsilon)}{2 \epsilon}
    $$

    donde usamos  $p$  de forma genérica para referirnos a cada elemento
     $\hookrightarrow w_{1\{i,j\}}, b_{1\{j\}}, w_{2\{i,j\}}, b_{2\{j\}}$ 
    """

    dW1 = np.zeros_like(self.W1)
    db1 = np.zeros_like(self.b1)
    dW2 = np.zeros_like(self.W2)
    db2 = np.zeros_like(self.b2)

    if dropout:
        temp_w1 = self.W1.copy()
        self.W1 = self.dropout(self.W1, dropout_prob)
        temp_b1 = self.b1.copy()
        self.b1 = self.dropout(self.b1, dropout_prob)
        temp_w2 = self.W2.copy()
        self.W2 = self.dropout(self.W2, dropout_prob)
        temp_b2 = self.b2.copy()
        self.b2 = self.dropout(self.b2, dropout_prob)

    for i in range(self.W1.shape[0]):
        for j in range(self.W1.shape[1]):
            if dropout:
                if self.W1[i, j] == 0 and temp_w1[i, j] != 0:
                    continue

```

```

        self.W1[i, j] += eps
        loss1 = self.loss(x, y)
        self.W1[i, j] -= 2 * eps
        loss2 = self.loss(x, y)
        dW1[i, j] = (loss1 - loss2) / (2 * eps)
        self.W1[i, j] += eps

    for i in range(self.b1.shape[0]):
        for j in range(self.b1.shape[1]):
            if dropout:
                if self.b1[i, j] == 0 and temp_b1[i, j] != 0:
                    continue
                self.b1[i, j] += eps
                loss1 = self.loss(x, y)
                self.b1[i, j] -= 2 * eps
                loss2 = self.loss(x, y)
                db1[i, j] = (loss1 - loss2) / (2 * eps)
                self.b1[i, j] += eps

    for i in range(self.W2.shape[0]):
        for j in range(self.W2.shape[1]):
            if dropout:
                if self.W2[i, j] == 0 and temp_w2[i, j] != 0:
                    continue
                self.W2[i, j] += eps
                loss1 = self.loss(x, y)
                self.W2[i, j] -= 2 * eps
                loss2 = self.loss(x, y)
                dW2[i, j] = (loss1 - loss2) / (2 * eps)
                self.W2[i, j] += eps

    for i in range(self.b2.shape[0]):
        for j in range(self.b2.shape[1]):
            if dropout:
                if self.b2[i, j] == 0 and temp_b2[i, j] != 0:
                    continue
                self.b2[i, j] += eps
                loss1 = self.loss(x, y)
                self.b2[i, j] -= 2 * eps
                loss2 = self.loss(x, y)
                db2[i, j] = (loss1 - loss2) / (2 * eps)
                self.b2[i, j] += eps

    if dropout:
        self.W1 = temp_w1
        self.b1 = temp_b1
        self.W2 = temp_w2

```

```

        self.b2 = temp_b2

    return dW1, db1, dW2, db2

def dropout(self, x: ndarray, p: float) -> ndarray:
    """
    Aplica dropout a la capa de entrada.
    x: ndarray de entrada
    p: probabilidad de dropout
    """
    mask = np.random.binomial(1, p, size=x.shape)
    return x * mask

def update_weights(
    self,
    lr: float,
    dW1: ndarray,
    db1: ndarray,
    dW2: ndarray,
    db2: ndarray,
):
    """
    Actualiza los pesos de la red neuronal usando
    numerical gradient descent.
    """

    self.W1 -= lr * dW1
    self.b1 -= lr * db1
    self.W2 -= lr * dW2
    self.b2 -= lr * db2

def loss(self, x: ndarray, y: ndarray, dropout: bool = False) -> float:
    """
    Calcula el error cuadrático medio.
    """

    return np.power((self.forward(x, dropout=dropout) - y), 2).mean(axis=0)
↪ / 2

def record_metrics(self, x: ndarray, y: ndarray, x_test: ndarray, y_test:
↪ ndarray):

    self.training_loss_acum.append(self.loss(x, y))
    self.test_loss_acum.append(self.loss(x_test, y_test))

def fit(
    self,

```

```

x: ndarray,
y: ndarray,
x_test: ndarray,
y_test: ndarray,
lr: float = 0.01,
epochs: int = 1000,
eps: float = 1e-3,
dropout: bool = False,
dropout_prob: float = 0.5,
custom_activation: Callable[[ndarray], ndarray] = None,
verbose: bool = False,
) -> List[float]:
    """
    Entrena la red neuronal usando gradient descent.
    """

    if custom_activation:
        self.activation = custom_activation

    self.train_loss_acum = []
    self.test_loss_acum = []

    self.lr = lr
    self.epochs = epochs
    self.eps = eps
    self.with_dropout = dropout
    self.dropout_prob = dropout_prob

    for _ in tqdm(range(epochs)) if verbose else range(epochs):
        self.record_metrics(x, y, x_test, y_test)

        dW1, db1, dW2, db2 = self.numerical_gradient(x, y, eps,
↪ dropout=dropout, dropout_prob=dropout_prob)

        self.update_weights(lr, dW1, db1, dW2, db2)

        self.activation = self.sigmoid

    return self.train_loss_acum

def predict(self, x: ndarray) -> ndarray:
    """
    Infiere la salida de la red neuronal.
    """
    return self.forward(x)

def get_weights(self) -> Tuple[ndarray, ndarray, ndarray, ndarray]:

```

```

    """
    Devuelve los pesos de la red neuronal.
    """

    return self.W1, self.b1, self.W2, self.b2

def get_training_loss(self) -> List[float]:
    """
    Devuelve el error cuadrático medio acumulado.
    """

    return self.training_loss_acum

def get_test_loss(self) -> List[float]:

    return self.test_loss_acum

def plot_loss(self, ax: plt.Axes = None) -> plt.Axes:

    if ax is None:
        _, ax = plt.subplots()

    ax.plot(
        range(len(self.training_loss_acum)),
        self.training_loss_acum,
        label="Training Loss",
    )
    ax.plot(range(len(self.test_loss_acum)), self.test_loss_acum,
    ↪label="Test Loss")
    ax.legend()
    ax.set_xlabel("Epoch")
    ax.set_ylabel("Loss")
    ax.set_title("Loss vs Epoch, lr = {}, $\epsilon$ = {}".format(self.lr,
    ↪self.eps))
    # If dropout is used, add it to the title
    if self.with_dropout:
        ax.set_title(
            "Loss vs Epoch, lr = {}, $\epsilon$ = {}, dropout_prob = {}".
    ↪format(
                self.lr, self.eps, self.dropout_prob
            )
        )

    return ax

def mse(self, y_true: ndarray, x_test: ndarray) -> float:
    """

```



```

Calcula el error cuadrático medio.
"""

y_pred = self.predict(x_test)
return np.mean(np.power(y_true - y_pred, 2))

def plot_network_graph(self, ax: plt.Axes):
    """
    Visualiza la red neuronal y sus conexiones
    """

    G = nx.DiGraph()
    G.add_nodes_from(["x1", "x2", "x3", "x4", "x5"])

    G.add_nodes_from(["z1_x1", "z1_x2", "z1_x3", "z1_x4", "z1_x5"])

    G.add_nodes_from(["y"])

    # Add edges with weights
    for i in range(5):
        for j in range(5):
            G.add_edge(f"x{i+1}", f"z1_x{j+1}", weight=self.W1[i, j])
            G.add_edge(f"z1_x{j+1}", f"y", weight=1)

    pos = {
        "x1": (0, 0),
        "x2": (0, 1),
        "x3": (0, 2),
        "x4": (0, 3),
        "x5": (0, 4),
        "z1_x1": (1, 0),
        "z1_x2": (1, 1),
        "z1_x3": (1, 2),
        "z1_x4": (1, 3),
        "z1_x5": (1, 4),
        "y": (2, 2),
    }

    # Edge labels with weights


    nx.draw_networkx_nodes(G, pos, cmap=plt.get_cmap("jet"),
↪node_size=2000, ax=ax)
    nx.draw_networkx_labels(G, pos, ax=ax)
    nx.draw_networkx_edges(
        G, pos, edgelist=G.edges(), edge_color="k", arrows=True, ax=ax
    )

```

```
[ ]: nn = NeuralNet()
fig, ax = plt.subplots(1, 1, figsize=(15, 5))
#no border
_ = fig.subplots_adjust(left=0, right=1, bottom=0, top=1)
_ = ax.axis('off')
_ = nn.plot_network_graph(ax)
```

/Users/nacho/opt/anaconda3/envs/coding/lib/python3.8/site-packages/networkx/drawing/nx_pyplot.py:433: UserWarning: No data for colormapping provided via 'c'. Parameters 'cmap' will be ignored

```
node_collection = ax.scatter(
```




1.5 Experimentación Manual

```
[ ]: nn = NeuralNet()
nn.fit(X_train, y_train, x_test=X_test, y_test=y_test, epochs=1000, lr=0.001,
      verbose=True)
display(pd.Series(
    [
        y_test[0][0],
        nn.predict(X_test[0].reshape(1, -1))[0][0],
        nn.get_training_loss()[-1][0],
        nn.get_test_loss()[-1][0],
    ],
    index=["y_true", "y_pred", "train_loss", "test_loss"]
))
nn.plot_loss()
```

```
100%|      | 1000/1000 [00:04<00:00, 222.64it/s]
```

```
y_true      27.30000000000000071054
y_pred      38.02465718134914141046
train_loss   97.12082915121834503225
test_loss    77.15179309729562362463
dtype: float64
```

```
[ ]: <Axes: title={'center': 'Loss vs Epoch, lr = 0.001,  $\epsilon$  = 0.001'},
      xlabel='Epoch', ylabel='Loss'>
```




```
[ ]: nn = NeuralNet()
nn.fit(X_train, y_train, x_test=X_test, y_test=y_test, epochs=1000, lr=0.001,
      custom_activation=nn.tanh, verbose=True)
display(pd.Series(
    [
        y_test[0][0],
        nn.predict(X_test[0].reshape(1, -1))[0][0],
        nn.get_training_loss()[-1][0],
        nn.get_test_loss()[-1][0],
    ],
    index=["y_true", "y_pred", "train_loss", "test_loss"]
))
nn.plot_loss()
```

```
100%|      | 1000/1000 [00:04<00:00, 248.22it/s]
```

```
y_true      27.3000000000000071054
y_pred      38.02122515605474006861
```

```
train_loss    97.12112880050391083842
test_loss     77.15022996002481647793
dtype: float64
```


```
[ ]: <Axes: title={'center': 'Loss vs Epoch, lr = 0.001,  $\epsilon$  = 0.001'},
      xlabel='Epoch', ylabel='Loss'>
```



```
[ ]: nn = NeuralNet()
nn.fit(X_train, y_train, x_test=X_test, y_test=y_test, epochs=1000, lr=0.001,
      ↪eps=1, verbose=True)
display(pd.Series(
    [
        y_test[0][0],
        nn.predict(X_test[0].reshape(1, -1))[0][0],
        nn.get_training_loss()[-1][0],
        nn.get_test_loss()[-1][0],
    ],
    index=["y_true", "y_pred", "train_loss", "test_loss"]
))
_ = nn.plot_loss()
```

100%| | 1000/1000 [00:04<00:00, 220.55it/s]

```
y_true      27.3000000000000071054
y_pred      27.87010821285398520786
train_loss   52.58049000028180586241
test_loss    35.92346049562767973384
dtype: float64
```

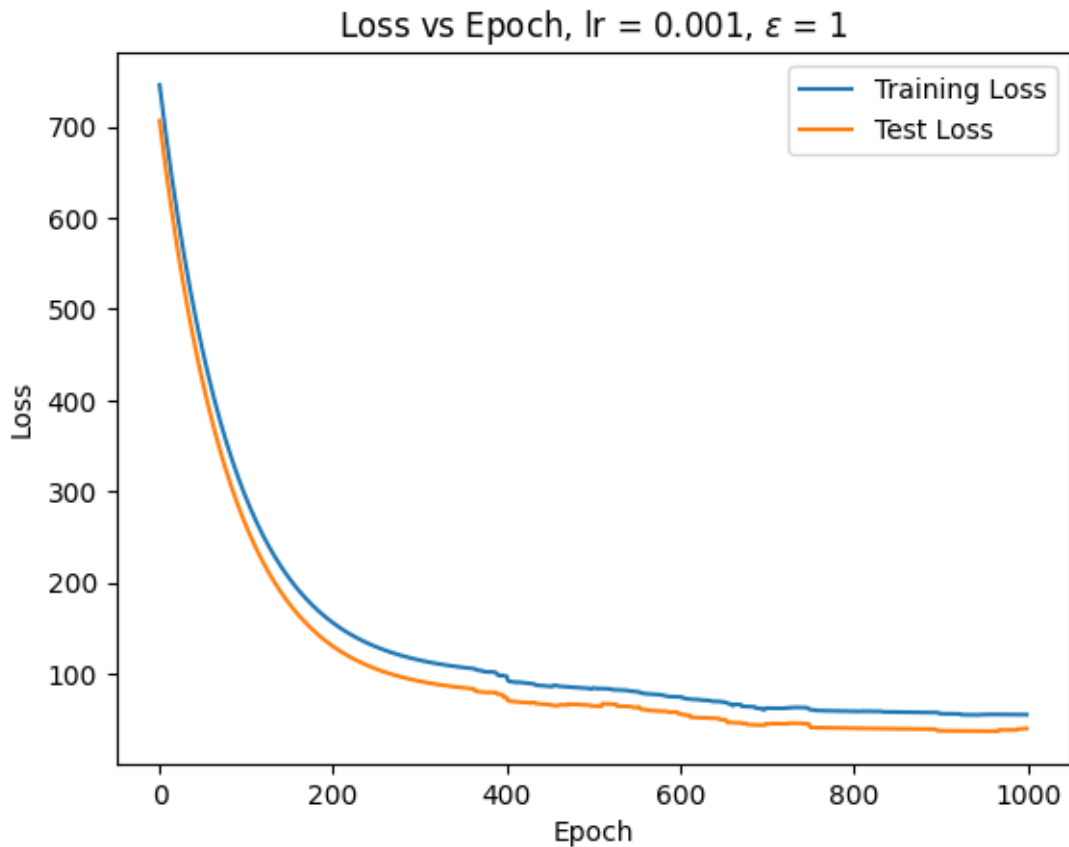


```
[ ]: nn = NeuralNet()
nn.fit(X_train, y_train, x_test=X_test, y_test=y_test, epochs=1000, lr=0.001,
      eps=1, custom_activation=nn.tanh, verbose=True)
display(pd.Series(
    [
        y_test[0][0],
        nn.predict(X_test[0].reshape(1, -1))[0][0],
        nn.get_training_loss()[-1][0],
        nn.get_test_loss()[-1][0],
    ],
    index=["y_true", "y_pred", "train_loss", "test_loss"]
))
```

```
_ = nn.plot_loss()
```

```
100%|      | 1000/1000 [00:04<00:00, 233.20it/s]
```

```
y_true      27.30000000000000071054
y_pred      35.42288054345024050917
train_loss   54.96998128619167545139
test_loss    39.69249387913455251464
dtype: float64
```



```
[ ]: nn_dropout = NeuralNet()
nn_dropout.fit(X_train, y_train, x_test=X_test, y_test=y_test, epochs=1000,
               lr=0.001, eps=1, dropout=True, dropout_prob=0.5, verbose=True)
display(pd.Series(
    [
        y_test[0][0],
        nn_dropout.predict(X_test[0].reshape(1, -1))[0][0],
        nn_dropout.get_training_loss()[-1][0],
        nn_dropout.get_test_loss()[-1][0],
        min(nn_dropout.get_training_loss())[0],
```

```

        min(nn_dropout.get_test_loss())[0],
    ],
    index=["y_true", "y_pred", "train_loss", "test_loss", "min_train_loss",
↪ "min_test_loss"]
))
_ = nn_dropout.plot_loss()


```

100%| | 1000/1000 [00:02<00:00, 399.85it/s]

```

y_true      27.30000000000000071054
y_pred      54.37790898914906279060
train_loss  228.93978918457611371196
test_loss   217.77674017536915584969
min_train_loss  97.11762557642698823201
min_test_loss  77.04854816912815351770
dtype: float64

```



```
[ ]: nn_dropout = NeuralNet()
```

```

nn_dropout.fit(X_train, y_train, x_test=X_test, y_test=y_test, epochs=1000,
↳lr=0.001, eps=1, dropout=True, dropout_prob=0.5, custom_activation=nn.tanh,
↳verbose=True)
display(pd.Series(
    [
        y_test[0][0],
        nn_dropout.predict(X_test[0].reshape(1, -1))[0][0],
        nn_dropout.get_training_loss()[-1][0],
        nn_dropout.get_test_loss()[-1][0],
        min(nn_dropout.get_training_loss())[0],
        min(nn_dropout.get_test_loss())[0],
    ],
    index=["y_true", "y_pred", "train_loss", "test_loss", "min_train_loss",
↳"min_test_loss"]
))
_ = nn_dropout.plot_loss()

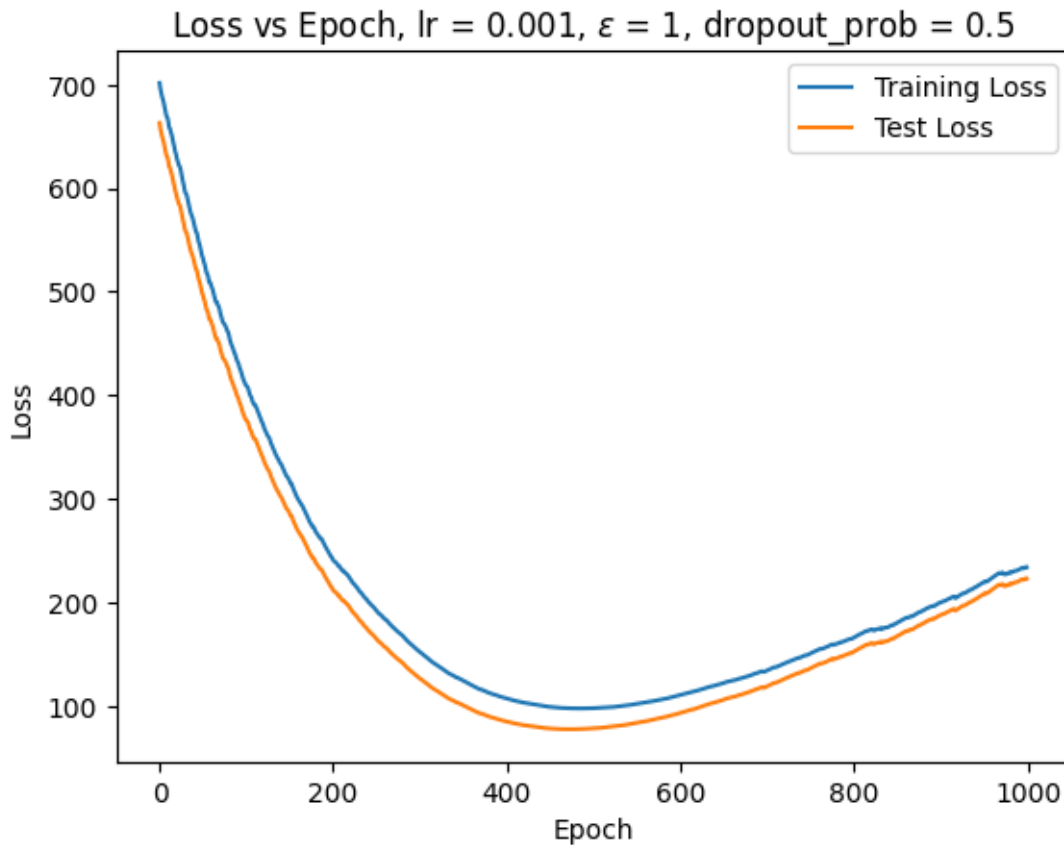
```

100%| | 1000/1000 [00:02<00:00, 461.19it/s]

```

y_true      27.30000000000000071054
y_pred      54.65448618360453281184
train_loss  233.43453540371953636168
test_loss   222.41958081889217169191
min_train_loss  97.11738405032269838557
min_test_loss  77.04853250241073681082
dtype: float64

```

```
[ ]: nn_dropout = NeuralNet()
nn_dropout.fit(X_train, y_train, x_test=X_test, y_test=y_test, epochs=10000,
↳lr=0.001, eps=1, dropout=True, dropout_prob=0.8, verbose=True)
```

```
100%|      | 10000/10000 [00:38<00:00, 259.49it/s]
```

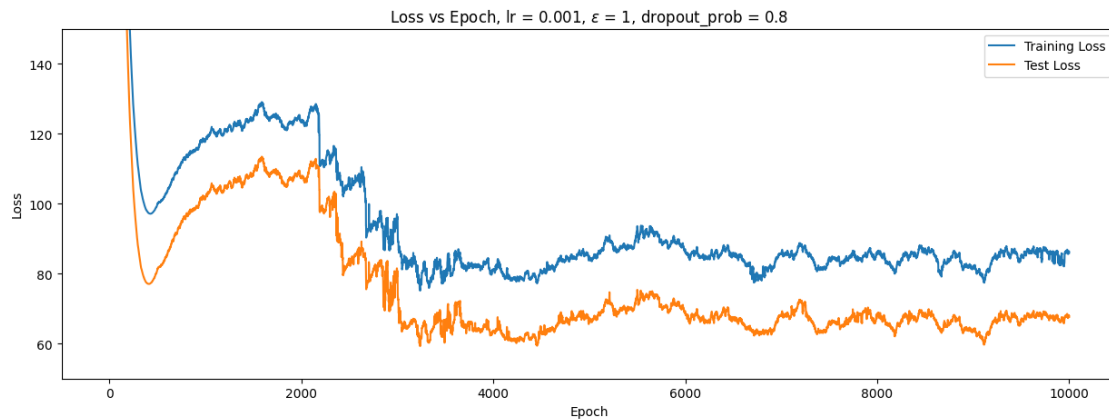
```
[ ]: []
```

```
[ ]: display(pd.Series(
    [
        y_test[0][0],
        nn_dropout.predict(X_test[0].reshape(1, -1))[0][0],
        nn_dropout.get_training_loss()[-1][0],
        nn_dropout.get_test_loss()[-1][0],
    ],
    index=["y_true", "y_pred", "train_loss", "test_loss"]
))
```

```
y_true      27.30000000000000071054
y_pred      33.30453404740335798806
train_loss   85.91754731507950282321
```


```
test_loss      67.61037497622338321435
dtype: float64
```

```
[ ]: fig, ax = plt.subplots(1, 1, figsize=(15, 5))
     ax.set_ylim(50, 150)
     _ = nn_dropout.plot_loss(ax=ax)
```



```
[ ]: nn_400_dropout = NeuralNet()
nn_400_dropout.fit(X_train, y_train, x_test=X_test, y_test=y_test, epochs=400,
    ↪lr=0.001, eps=0.1, dropout=True, dropout_prob=0.8)
display(pd.Series(
    [
        y_test[0][0],
        nn_400_dropout.predict(X_test[0].reshape(1, -1))[0][0],
        nn_400_dropout.get_training_loss()[-1][0],
        nn_400_dropout.get_test_loss()[-1][0],
    ],
    index=["y_true", "y_pred", "train_loss", "test_loss"]
))
_ = nn_400_dropout.plot_loss()
```

```
y_true      27.30000000000000071054
y_pred      37.03763934176618022320
train_loss   97.69831781165335371497
test_loss    77.19354891780324123829
dtype: float64
```



```
[ ]: hyper_nn = NeuralNet()
hyper_nn.fit(X_train, y_train, x_test=X_test, y_test=y_test, epochs=50000, lr=0.
↳001, eps=1, dropout=True, dropout_prob=0.8)
```

```
/var/folders/5j/7b3w3dw95qdc1l87vxc3yn300000gn/T/ipykernel_1661/2190381686.py:28
: RuntimeWarning: overflow encountered in exp
return 1 / (1 + np.exp(-x))
```

```
[ ]: []
```

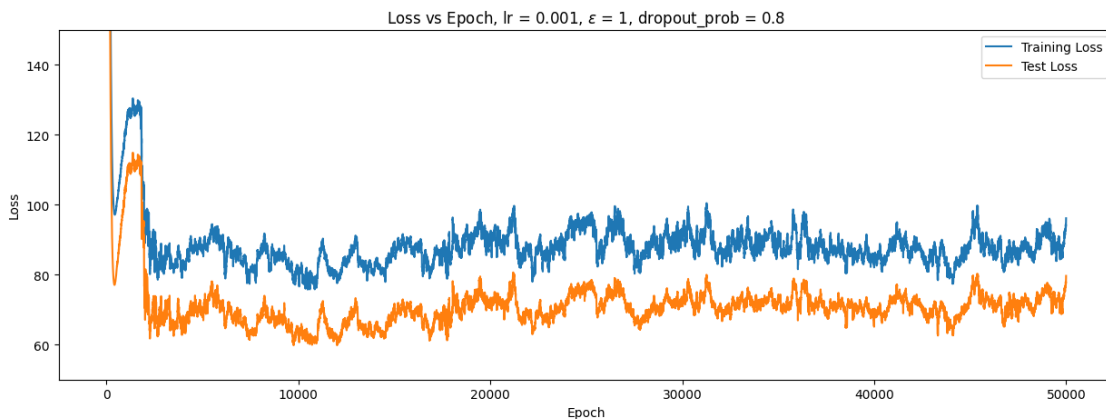
```
[ ]: display(pd.Series(
    [
        y_test[0][0],
        hyper_nn.predict(X_test[0].reshape(1, -1))[0][0],
        hyper_nn.get_training_loss()[-1][0],
        hyper_nn.get_test_loss()[-1][0],
    ],
    index=["y_true", "y_pred", "train_loss", "test_loss"]
))
```

```
/var/folders/5j/7b3w3dw95qdc1l87vxc3yn300000gn/T/ipykernel_1661/2190381686.py:28
```

```
: RuntimeWarning: overflow encountered in exp
    return 1 / (1 + np.exp(-x))
```

```
y_true      27.30000000000000071054
y_pred      28.12064270715838176784
train_loss   96.12452944561505319143
test_loss    79.64967781583217742991
dtype: float64
```

```
[ ]: fig, ax = plt.subplots(1, 1, figsize=(15, 5))
    ax.set_ylim(50, 150)
    hyper_nn.plot_loss(ax=ax)
```



```
[ ]: # Epoch en el que se obtuvo la menor Loss

m = min(hyper_nn.get_training_loss(), key=lambda x: x[0])
e = hyper_nn.get_training_loss().index(m)

print(f"Lowest training loss: {m[0]:.2f} at epoch {e}")
```

Lowest training loss: 75.74 at epoch 10480

```
[ ]: hyper_nn_ = NeuralNet()
    hyper_nn_.fit(X_train, y_train, x_test=X_test, y_test=y_test, epochs=e, lr=0.
        ↪001, eps=1, dropout=True, dropout_prob=0.8)
```

1.6 Mascara sobre los pesos: Dropout


Otras dos cuestiones que probamos fueron cambiar la función de activación, y aplicar una mascara a los pesos de entrenamiento de la red con cierta probabilidad en cada epoch.

Para la función de activación probamos con la función sigmoide como indicaba la consigna, y ademas con la funciones ReLU y Tanh.

$$\text{ReLU}(x) = \max(0, x) \quad \text{Tanh}(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$


```
[ ]: nn = NeuralNet()
loss = nn.fit(X_train, y_train, x_test=X_test, y_test=y_test, lr=0.003,
             epochs=400, dropout=True)
display(pd.Series(
    [
        y_test[0][0],
        nn.predict(X_test[0].reshape(1, -1))[0][0],
        nn.get_training_loss()[-1][0]
    ],
    index=["y_true", "y_pred", "train_loss"]
))
nn.plot_loss()
```

```
y_true      27.30000000000000071054
y_pred      57.63734020134181434969
train_loss   287.87181630280230137942
dtype: float64
```



```
[ ]: nn = NeuralNet()
loss = nn.fit(X_train, y_train, x_test=X_test, y_test=y_test, lr=0.1, epochs=6,
↳ dropout=True)
display(pd.Series(
    [
        y_test[0][0],
        nn.predict(X_test[0].reshape(1, -1))[0][0],
        nn.get_training_loss()[-1][0]
    ],
    index=["y_true", "y_pred", "train_loss"]
))
nn.plot_loss()
```

```
y_true      27.3000000000000071054
y_pred      46.12380957676615622631
train_loss   107.23579872377744948153
dtype: float64
```



```
[ ]: nn = NeuralNet()
```

```

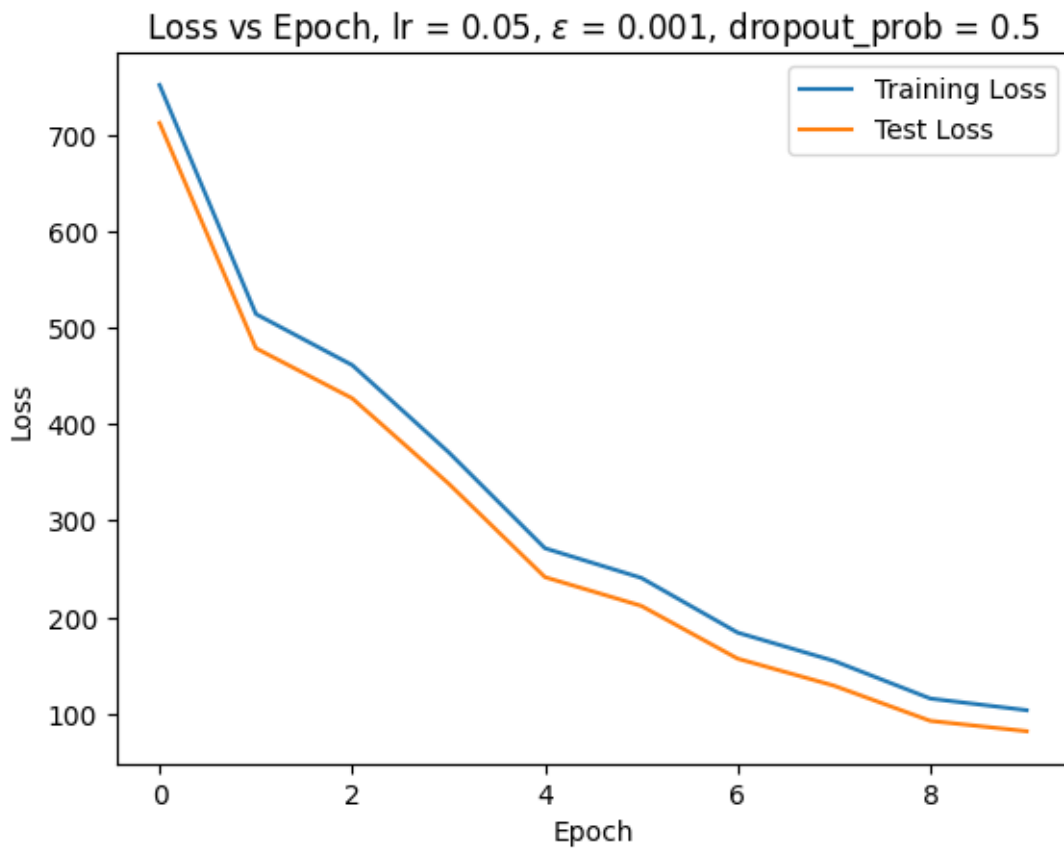
loss = nn.fit(X_train, y_train, x_test=X_test, y_test=y_test, lr=0.05,
epochs=10, dropout=True)
display(pd.Series(
    [
        y_test[0][0],
        nn.predict(X_test[0].reshape(1, -1))[0][0],
        nn.get_training_loss()[-1][0]
    ],
    index=["y_true", "y_pred", "train_loss"]
))
nn.plot_loss()

```

```

y_true      27.30000000000000071054
y_pred      38.24117112246570826528
train_loss   103.20305378438199284119
dtype: float64

```



1.7 Otras funciones de activación


Para la función de activación probamos con la función sigmoide como indicaba la consigna, y además con las funciones ReLU, Leaky Relu y Tanh.

$$\text{ReLU}(x) = \max(0, x) \quad \text{LeakyReLU}(x) = \max(0.01x, x) \quad \text{Tanh}(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

1.7.1 ReLu y Leaky ReLu

Con las función ReLU y Leaky ReLU no logramos convergencia.


```
[ ]: nn_relu = NeuralNet()  
nn_relu.fit(X_train, y_train, x_test=X_test, y_test=y_test, lr=0.003,   
           epochs=400, custom_activation=nn.relu)  
nn_relu.plot_loss()
```



```
[ ]: nn_leaky = NeuralNet()  
nn_leaky.fit(X_train, y_train, x_test=X_test, y_test=y_test, lr=0.003,   
            epochs=400, custom_activation=nn.leaky_relu)  
nn_leaky.plot_loss()
```


Sin Embargo para la función Tanh si vimos convergencia en nuestras pruebas.

```
[ ]: nn_tanh = NeuralNet()
nn_tanh.fit(X_train, y_train, x_test=X_test, y_test=y_test, lr=0.003,
↳ epochs=400, custom_activation=nn.tanh)
nn_tanh.plot_loss()
```



Aca una comparación entre las funciones de activación sigmoide y tanh. En este caso la función de activación tanh hasta converge mas rápido y a una loss menor que la sigmoide.

```
[ ]: fig, ax = plt.subplots(figsize=(15, 5))

nn = NeuralNet()
nn.fit(X_train, y_train, x_test=X_test, y_test=y_test, lr=0.003, epochs=150)

nn_tanh = NeuralNet()
nn_tanh.fit(X_train, y_train, x_test=X_test, y_test=y_test, lr=0.003,
↳ epochs=150, custom_activation=nn.tanh)

ax.plot(nn.get_training_loss(), label="Sigmoid Training Loss")
```


```

ax.plot(nn.get_test_loss(), label="Sigmoid Test Loss")

ax.plot(nn_tanh.get_training_loss(), label="Tanh Training Loss")
ax.plot(nn_tanh.get_test_loss(), label="Tanh Test Loss")
ax.legend()

```

```
[ ]: <matplotlib.legend.Legend at 0x7faeb6ae5d30>
```



1.8 Distintos valores de Learning Rate para distintos Epochs

```
[ ]: # Brute force para encontrar el mejor número de epochs y learning rate
```

```

lrs = np.linspace(0.001, 0.5, 10).round(3)
eps = list(range(1, 1000, 50))

train_losses = {}
test_losses = {}

for lr, ep in tqdm(list(itertools.product(lrs, eps))):
    nn = NeuralNet()
    nn.fit(X_train, y_train, x_test=X_test, y_test=y_test, lr=lr,
    ↪ epochs=int(ep))

    train_losses[(lr, int(ep))] = nn.get_training_loss()[-1][0]
    test_losses[(lr, int(ep))] = nn.get_test_loss()[-1][0]

```

```
100%|      | 200/200 [06:15<00:00, 1.88s/it]
```

```

[ ]: train_losses_arr = []
test_losses_arr = []

for (lr, ep), loss in train_losses.items():
    train_losses_arr.append((lr, int(ep), loss))

```

```
[ ]: lr epochs \
```

0	0.0010000000000000000002	1
1	0.0010000000000000000002	51
2	0.0010000000000000000002	101
3	0.0010000000000000000002	151
4	0.0010000000000000000002	201
..
195	0.5000000000000000000000	751
196	0.5000000000000000000000	801
197	0.5000000000000000000000	851
198	0.5000000000000000000000	901
199	0.5000000000000000000000	951

```
loss_train \
```

0	712.35544408202019894816
1	438.10390531169031191894
2	277.70827155145241249556
3	204.45153606491405184897
4	152.95659639606813584578
..	...
195	10,064,873,608,080,629,039,315,616,564,842,070,...
196	38,883,928,850,905,498,030,357,258,510,008,320...
197	1,964,394,829,934,937,164,113,258,741,760.00000...
198	161,643,292,004,524,220,750,446,944,777,003,859...
199	5,408,037,911,253,026,769,260,384,748,568,576.0...

```
loss_test
```

0	673.50730859418263207772
1	404.09167713681171107964
2	247.53186235583663687976
3	176.62372830443592874872
4	127.33201401750525860734
..	...
195	10,064,873,608,080,626,678,132,375,130,019,463,...

```

196 38,883,928,850,905,493,418,671,240,082,620,416...
197 1,964,394,829,934,938,571,488,142,295,040.00000...
198 161,643,292,004,524,220,750,446,944,777,003,859...
199 5,408,037,911,253,026,769,260,384,748,568,576.0...

```

[200 rows x 4 columns]

```
[ ]: losses_df.sort_values(by='loss_test', ascending=True).head(10)
```

```

[ ]:
      lr  epochs  loss_train \
14  0.001000000000000000002    701 97.24997849157063001257
13  0.001000000000000000002    651 97.36517063828993912011
15  0.001000000000000000002    751 97.18700042019692375561
16  0.001000000000000000002    801 97.15765470153172600476
17  0.001000000000000000002    851 97.13886442836047763194
12  0.001000000000000000002    601 97.53729094639008678769
18  0.001000000000000000002    901 97.12979841319094020946
19  0.001000000000000000002    951 97.12443781084131444459
148 0.3890000000000000001243    401 97.11721857464502249968
182 0.500000000000000000000    101 97.11721191618177329019

      loss_test
14  77.04882381592535978143
13  77.06209465460953822458
15  77.06229356767643423609
16  77.08106987569057366727
17  77.10344679716750704301
12  77.11963007880184761689
18  77.12105310560889392946
19  77.13643326268517341759
148 77.19205646245166008157
182 77.19363039787369018541

```

```
[ ]: losses_df.sort_values(by='loss_train', ascending=True).head(10)
```

```

[ ]:
      lr  epochs  loss_train \
81  0.223000000000000000377     51 97.11721158982109614044
98  0.223000000000000000377    901 97.11721158982109614044
118 0.2780000000000000002487    901 97.11721158982109614044
26  0.056000000000000000117    301 97.11721158982109614044
95  0.223000000000000000377    751 97.11721158982109614044
45  0.112000000000000000233    251 97.11721158982109614044
42  0.112000000000000000233    101 97.11721158982109614044
44  0.112000000000000000233    201 97.11721158982109614044
28  0.056000000000000000117    401 97.11721158982109614044
71  0.167000000000000000955    551 97.11721158982109614044

```

```

            loss_test
81  77.19406594744638994143
98  77.19406594746409666641
118 77.19406594745940708435
26  77.19406594745862548734
95  77.19406594745851180051
45  77.19406594745805705315
42  77.19406594745773020350
44  77.19406594745751704068
28  77.19406594745743177555
71  77.19406594745981919914

```

Muchos de los valores de lr convergen a la misma loss minima $\sim(97.117)$, pero algunos lo hacen mas rapido que otros. En nuestra experimentación con el Learning Rate, obtuvimos que para $lr=0.05$ y $epoch=20$, la Loss alcanzaba $\sim(97.118)$.

En estos gráficos de calor podemos ver como la mayoría de los learning rates para cierto epoch logran converger a este mismo valor.


```

[ ]: losses_df_hm = losses_df.pivot(index='lr', columns='epochs', values='loss_test')

fig, ax = plt.subplots(figsize=(15, 10))

sns.heatmap(losses_df_hm, ax=ax, cmap='viridis')
ax.set_xlabel('Epochs')
ax.set_ylabel('Learning Rate')
ax.set_title('Test Loss Heatmap')
ax.invert_yaxis()
plt.show()

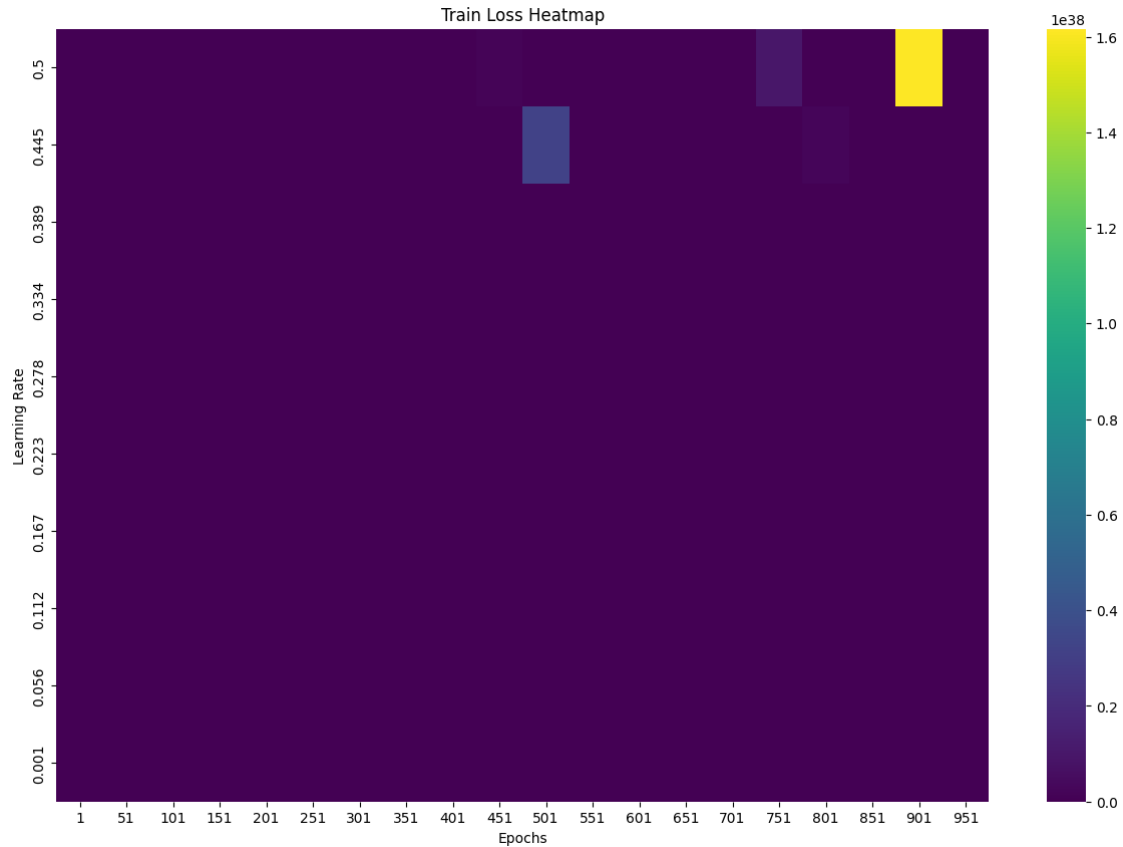
```



```
[ ]: losses_df_hm = losses_df.pivot(index='lr', columns='epochs',
    ↪ values='loss_train')

fig, ax = plt.subplots(figsize=(15, 10))


sns.heatmap(losses_df_hm, ax=ax, cmap='viridis')
ax.set_xlabel('Epochs')
ax.set_ylabel('Learning Rate')
ax.set_title('Train Loss Heatmap')
ax.invert_yaxis()
plt.show()
```



Para lograr obtener una menor Loss, probamos modificando el valor de Epsilon

```
[ ]: nn = NeuralNet()
loss = nn.fit(X_train, y_train, x_test=X_test, y_test=y_test, lr=0.01,
             epochs=100, eps=1)
display(pd.Series(
    [
        y_test[0][0],
        nn.predict(X_test[0].reshape(1, -1))[0][0],
        nn.get_training_loss()[-1][0]
    ],
    index=["y_true", "y_pred", "train_loss"]
))
nn.plot_loss()
```

```
y_true      27.30000000000000071054
y_pred      28.85025195570380773802
train_loss   53.80364135528142099929
dtype: float64
```



1.9 Combinaciones de Learning Rate y Epsilon

Buscamos ver como variando el Epsilon junto al Learning Rate para una cantidad fija arbitraria de epochs podemos encontrar un valor de lr que converja mas rápido sobre los modelos con las funciones de activación Tanh y Sigmoide, esta última con y sin Dropout de pesos. Además, vamos a ver como cambia la loss minima para cuando finalice el entrenamiento.

```
[ ]: # Brute force, best learning rate, epsilons, with and without dropout

lrs = np.linspace(0.001, 0.5, 10).round(3)
#epochs = list(range(0, 1050, 50))[1:]
epochs = [500] # Se podría probar con distintos epochs, idealmente valores
               ↪ mayores, los anteriores se pueden obtener como "early stops"
epsilons = list(map(lambda x: x/10000, range(0, 10005, 500)[1:]))

[ ]: len(epsilons), len(lrs), len(epochs)

[ ]: (20, 10, 1)

[ ]: len(list(itertools.product(lrs, epochs, epsilons)))
```



```
[ ]: 200
```

```
[ ]: losses = {  
    'default': {  
        'train': {},  
        'test': {}  
    },  
    'tanh': {  
        'train': {},  
        'test': {}  
    },  
    'dropout': {  
        'train': {},  
        'test': {}  
    },  
}
```

```
[ ]: for lr, ep, eps in tqdm(list(itertools.product(lrs, epochs, epsilons))):
```

```
    nn = NeuralNet()  
    nn.fit(X_train, y_train, x_test=X_test, y_test=y_test, lr=lr,  
↪ epochs=int(ep), eps=eps)  
    losses['default']['train'][lr, ep, eps] = nn.get_training_loss()  
    losses['default']['test'][lr, ep, eps] = nn.get_test_loss()  
  
    nn_drop = NeuralNet()  
    nn_drop.fit(X_train, y_train, x_test=X_test, y_test=y_test, lr=lr,  
↪ epochs=int(ep), eps=eps, dropout=True)  
    losses['dropout']['train'][lr, ep, eps] = nn_drop.get_training_loss()  
    losses['dropout']['test'][lr, ep, eps] = nn_drop.get_test_loss()  
  
    nn_tanh = NeuralNet()  
    nn_tanh.fit(X_train, y_train, x_test=X_test, y_test=y_test, lr=lr,  
↪ epochs=int(ep), eps=eps, custom_activation=nn_tanh.tanh)  
    losses['tanh']['train'][lr, ep, eps] = nn_tanh.get_training_loss()  
    losses['tanh']['test'][lr, ep, eps] = nn_tanh.get_test_loss()
```

```
1%|          | 2/200 [00:10<17:29, 5.30s/it]/var/folders/5j/7b3w3dw95qdc1l87v  
xc3yn300000gn/T/ipykernel_64997/343068776.py:30: RuntimeWarning: overflow  
encountered in exp
```

```
    return 1 / (1 + np.exp(-x))  
100%|         | 200/200 [17:34<00:00, 5.27s/it]
```

```
[ ]: fig, axs = plt.subplots(len(losses.keys()), 2, figsize=(15, 10))  
  
for i, (variation, train_test) in enumerate(losses.items()):  
    for (lr, ep, eps), loss in train_test['train'].items():
```

```

# De forma empírica observamos que con pesos aleatorios, la loss
↪ comienza al rededor de ~700
# Evitamos en el gráfico aquellos que superen en algún momento esa loss
↪ para limpiar
if any([l > 800 for l in loss]):
    continue
axs[i, 0].plot(loss, label=f'lr={lr}, ep={ep}, eps={eps}')

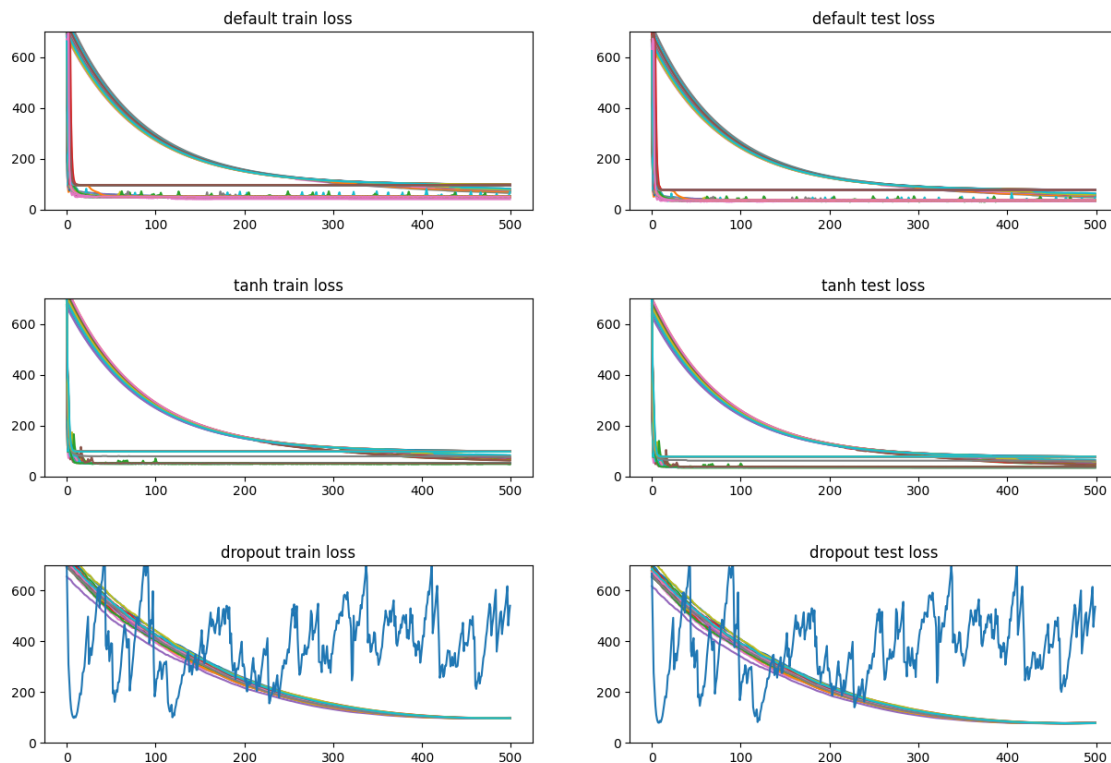
for (lr, ep, eps), loss in train_test['test'].items():
    if any([l > 800 for l in loss]):
        continue
    axs[i, 1].plot(loss, label=f'lr={lr}, ep={ep}, eps={eps}')

axs[i, 0].set_title(f'{variation} train loss')
axs[i, 1].set_title(f'{variation} test loss')

axs[i, 0].set_ylim(0, 700)
axs[i, 1].set_ylim(0, 700)

# Add space between subplots
fig.subplots_adjust(hspace=0.5)

```



Podemos observar como para ciertas combinaciones de Learning Rate y Epochs la convergencia es

mas rápida.

Para la función de activación ReLU observamos que la convergencia no ocurre directamente por lo que se dejó de evaluar, la función Tanh sin embargo si converge, pero no logra alcanzar la misma loss minima que la función sigmoide.

Para la mascara de pesos, probamos con una probabilidad de 0.5. En todos los casos la convergencia fue o mas lenta, y la loss minima alcanzada mayor que la obtenida sin aplicar la mascara, o directamente esporádico sin convergencia.

```
[ ]: # Large Default Training Loss Plot
fig, ax = plt.subplots(figsize=(20, 10))

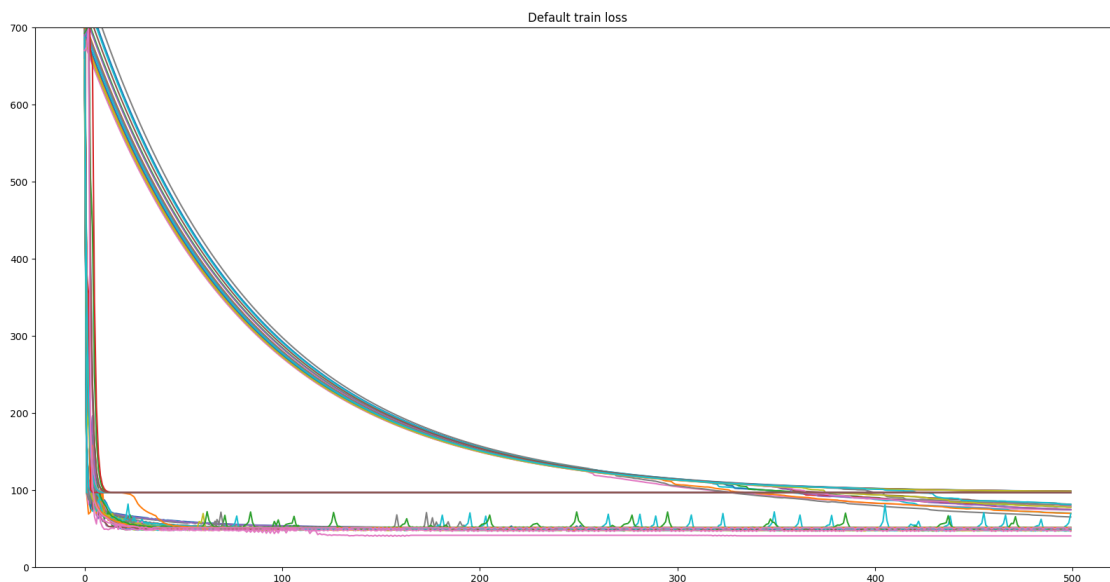
for (lr, ep, eps), loss in losses['default']['train'].items():

    if any([l > 800 for l in loss]):
        continue
    ax.plot(loss, label=f'lr={lr}, ep={ep}, eps={eps}')

ax.set_title(f'Default train loss')
#ax.legend()

ax.set_ylim(0, 700)
```

```
[ ]: (0.0, 700.0)
```



```
[ ]: # Large Default Training Loss Plot
fig, ax = plt.subplots(figsize=(20, 10))
```


```

for (lr, ep, eps), loss in losses['default']['train'].items():
    if any([l > 800 for l in loss]):
        continue
    ax.plot(loss, label=f'lr={lr}, ep={ep}, eps={eps}')

ax.set_title(f'Default train loss')
ax.set_ylim(40, 120)

```

[]: (40.0, 120.0)



```


[ ]: # Large Default Test Loss Plot
fig, ax = plt.subplots(figsize=(20, 10))

for (lr, ep, eps), loss in losses['default']['test'].items():
    if any([l > 800 for l in loss]):
        continue
    ax.plot(loss, label=f'lr={lr}, ep={ep}, eps={eps}')

ax.set_title(f'Default test loss')
ax.set_ylim(30, 80)

```

[]: (30.0, 80.0)



```
[ ]: # Large Default Training Loss Plot
fig, ax = plt.subplots(figsize=(20, 10))

for (lr, ep, eps), loss in losses['tanh']['train'].items():
    if any([l > 800 for l in loss]):
        continue
    ax.plot(loss, label=f'lr={lr}, ep={ep}, eps={eps}')

ax.set_title(f'tanh train loss')
ax.set_ylim(40, 120)
```

```
[ ]: (40.0, 120.0)
```



```
[ ]: # Filter out the best learning rate - epsilons combinations for the Default NN
# The 20 lowest training loss by epoch 200

# Sort by the loss at epoch 250

train_loss = losses['default']['train'].items()

train_loss = sorted(train_loss, key=lambda x: x[1][250])

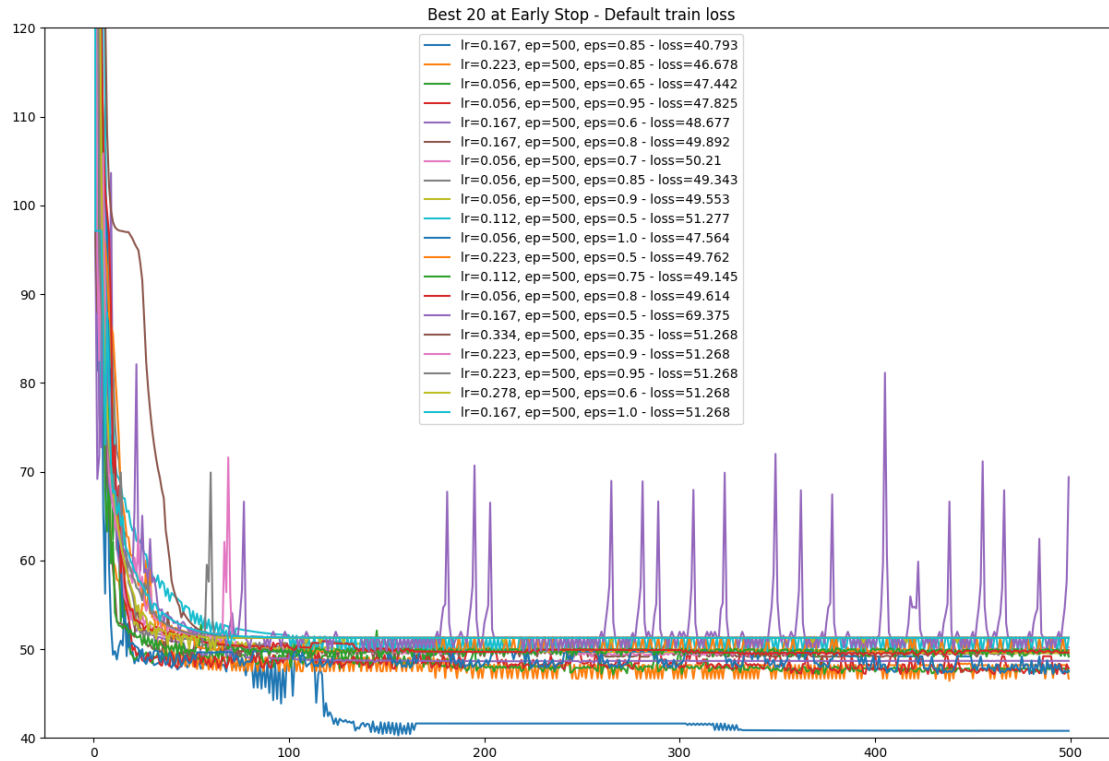
# Plot every train_loss[:20]

fig, ax = plt.subplots(figsize=(15, 10))

for (lr, ep, eps), loss in train_loss[:20]:
    ax.plot(loss, label=f'lr={lr}, ep={ep}, eps={eps} - loss={loss[-1].
        round(3)[0]}')

ax.set_title(f'Best 20 at Early Stop - Default train loss')
ax.set_ylim(40, 120)
ax.legend()
```

```
[ ]: <matplotlib.legend.Legend at 0x7faec6a079a0>
```



```
[ ]: # Filter out the best learning rate - epsilons combinations for the Default NN
# The 20 lowest training loss by epoch 200

# Sort by the loss at epoch 250

train_loss = losses['tanh']['train'].items()

train_loss = sorted(train_loss, key=lambda x: x[1][250])

# Plot every train_loss[:20]

fig, ax = plt.subplots(figsize=(15, 10))


for (lr, ep, eps), loss in train_loss[:20]:
    ax.plot(loss, label=f'lr={lr}, ep={ep}, eps={eps} - loss={loss[-1].
        round(3)[0]}')

(min_250_lr, min_250_ep, min_250_eps), min_250_loss = train_loss[0]

ax.set_title(f'Best 20 at Early Stop (Epoch 250) - Tanh train loss \n Best:␣
    lr={min_250_lr}, ep={min_250_ep}, eps={min_250_eps} - loss={min_250_loss[-1].
        round(3)[0]}')
```

```
ax.set_ylim(40, 120)
ax.legend()
```

```
[ ]: <matplotlib.legend.Legend at 0x7faec6e9d310>
```



```
[ ]: # Large Dropout Training Loss Plot
```

```
fig, ax = plt.subplots(figsize=(15, 10))

for (lr, ep, eps), loss in losses['dropout']['train'].items():

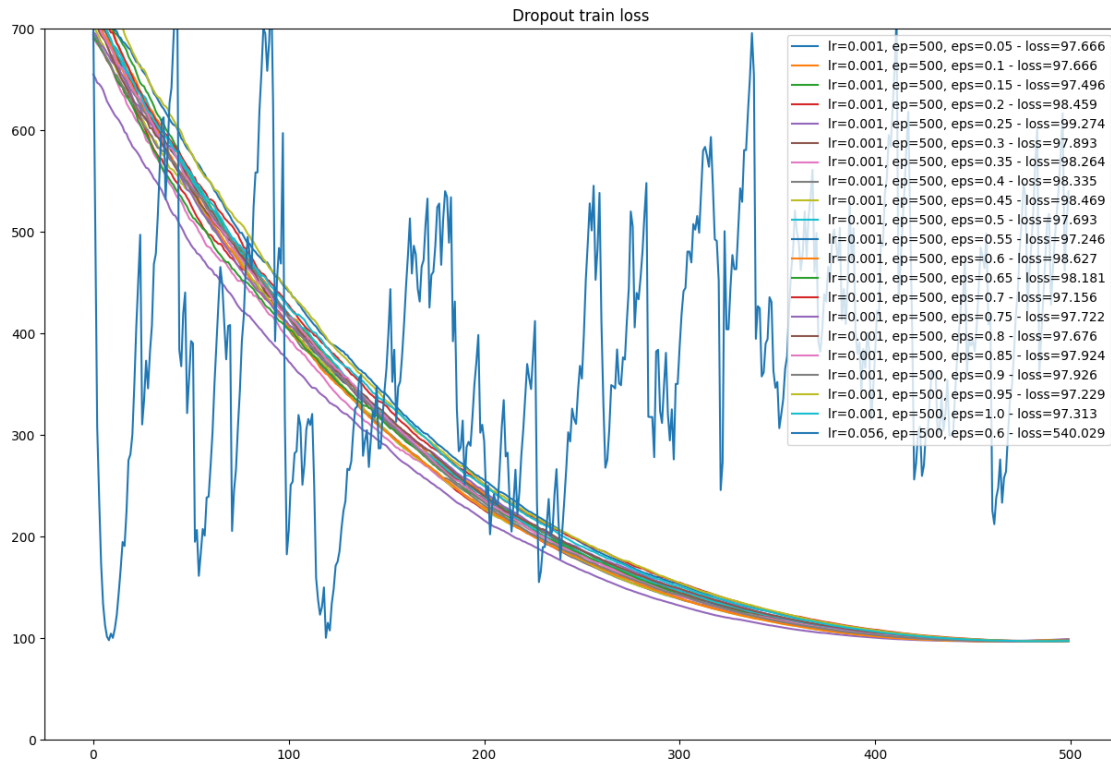
    # If loss has any value > 800, skip
    if any([l > 800 for l in loss]):
        continue

    ax.plot(loss, label=f'lr={lr}, ep={ep}, eps={eps} - loss={loss[-1].
    round(3)[0]}')

ax.set_title(f'Dropout train loss')
ax.set_ylim(0, 700)
ax.legend()
```



```
[ ]: <matplotlib.legend.Legend at 0x7faea3223520>
```



```
[ ]: # Min loss for default at epoch 100

(min_lr, min_ep, min_eps), min_train_loss = min(losses['default']['train'].
    ↪items(), key=lambda x: x[1][100])
min_test_loss = losses['default']['test'][(min_lr, min_ep, min_eps)]

fig, ax = plt.subplots(figsize=(15, 10))

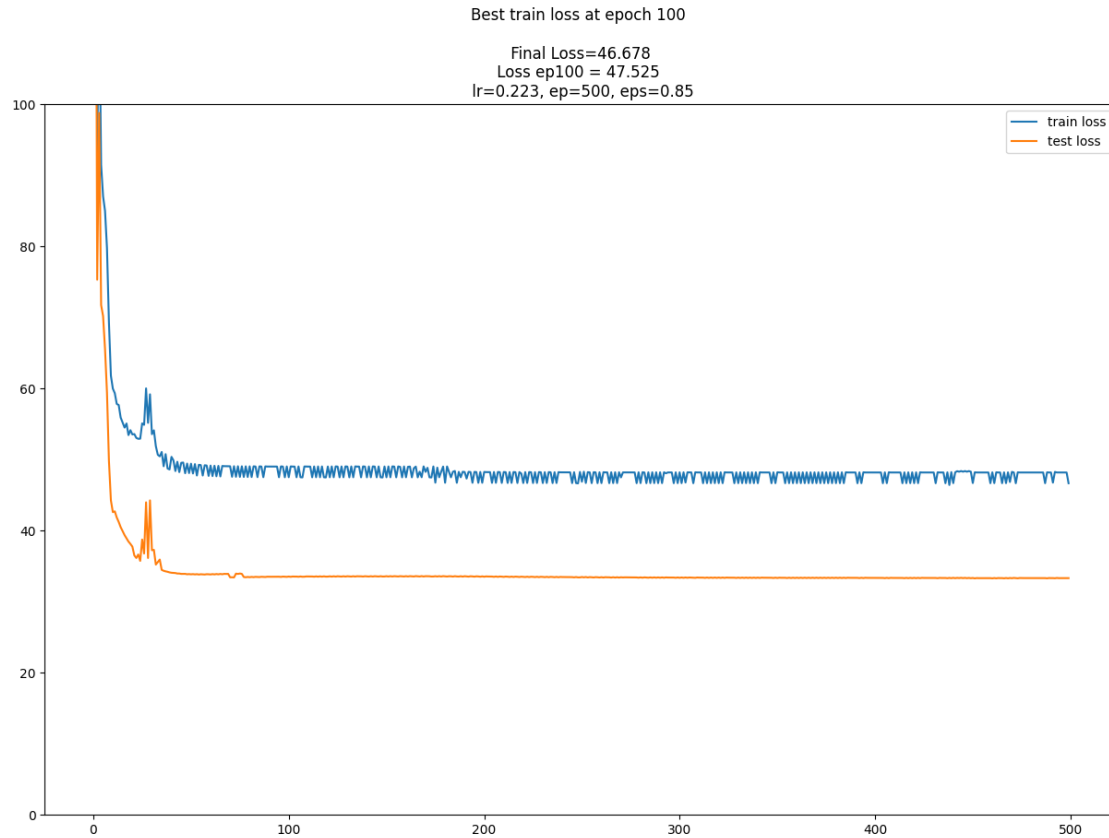
ax.plot(min_train_loss, label=f'train loss')
ax.plot(min_test_loss, label=f'test loss')

ax.set_title(f'Best train loss at epoch 100 \n\n Final Loss={min_train_loss[-1].
    ↪round(3)[0]} \n Loss ep100 = {min_train_loss[100].round(3)[0]} \n
    ↪lr={min_lr}, ep={min_ep}, eps={min_eps}')

ax.legend()

ax.set_ylim(0, 100)
```

```
[ ]: (0.0, 100.0)
```



```
[ ]: # Filter dropout losses, if any > 800

less_dropout_train_loss = {}
less_dropout_test_loss = {}

for (lr, ep, eps), loss in losses['dropout']['train'].items():
    if any([l > 800 for l in loss]):
        continue
    less_dropout_train_loss[(lr, ep, eps)] = loss
    less_dropout_test_loss[(lr, ep, eps)] = losses['dropout']['test'][(lr, ep,
↪eps)]
```

```
[ ]: # Min loss for default at epoch 500

(min_lr, min_ep, min_eps), min_train_loss = min(losses['default']['train'].
↪items(), key=lambda x: x[1][-1])
min_test_loss = losses['default']['test'][(min_lr, min_ep, min_eps)]

(min_tanh_lr, min_tanh_ep, min_tanh_eps), min_tanh_train_loss =
↪min(losses['tanh']['train'].items(), key=lambda x: x[1][-1])
```

```

min_tanh_test_loss = losses['tanh']['test'][(min_tanh_lr, min_tanh_ep,
↳min_tanh_eps)]

(min_dropout_lr, min_dropout_ep, min_dropout_eps), min_dropout_train_loss =
↳min(less_dropout_train_loss.items(), key=lambda x: x[1][-1])
min_dropout_test_loss = less_dropout_test_loss[(min_dropout_lr, min_dropout_ep,
↳min_dropout_eps)]

fig, ax = plt.subplots(figsize=(15, 10))

ax.plot(min_train_loss, label=f'Best final train loss')
ax.plot(min_test_loss, label=f'Best final test loss')

ax.plot(min_tanh_train_loss, label=f'Best final tanh train loss')
ax.plot(min_tanh_test_loss, label=f'Best final tanh test loss')

ax.plot(min_dropout_train_loss, label=f'Best final dropout train loss')
ax.plot(min_dropout_test_loss, label=f'Best final dropout test loss')

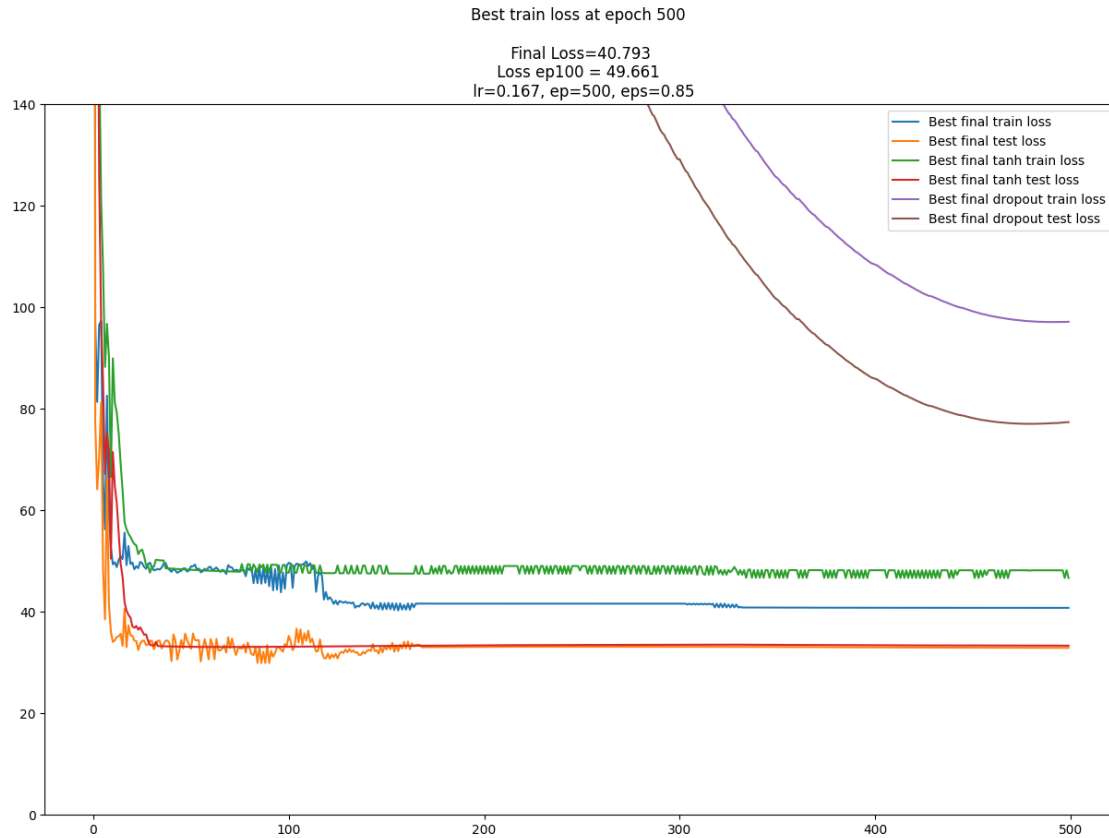
ax.set_title(f'Best train loss at epoch 500 \n\n Final Loss={min_train_loss[-1].
↳round(3)[0]} \n Loss ep100 = {min_train_loss[100].round(3)[0]} \n
↳lr={min_lr}, ep={min_ep}, eps={min_eps}')

ax.legend()

ax.set_ylim(0, 140)

```

```
[ ]: (0.0, 140.0)
```



```
[ ]: values_df = pd.DataFrame({
    'learning_rate': [min_lr, min_tanh_lr, min_dropout_lr],
    'epsilon': [min_eps, min_tanh_eps, min_dropout_eps],
    'epochs': [min_ep, min_tanh_ep, min_dropout_ep],
    'train_loss': [min_train_loss[-1].round(3)[0], min_tanh_train_loss[-1].
    ↪round(3)[0], min_dropout_train_loss[-1].round(3)[0]],
    'test_loss': [min_test_loss[-1].round(3)[0], min_tanh_test_loss[-1].
    ↪round(3)[0], min_dropout_test_loss[-1].round(3)[0]],
    'train_loss_ep100': [min_train_loss[100].round(3)[0], ↪
    ↪min_tanh_train_loss[100].round(3)[0], min_dropout_train_loss[100].
    ↪round(3)[0]],
    'test_loss_ep100': [min_test_loss[100].round(3)[0], min_tanh_test_loss[100].
    ↪round(3)[0], min_dropout_test_loss[100].round(3)[0]],
    }, index=['default', 'tanh', 'dropout'])

values_df
```

```
[ ]:
           learning_rate      epsilon  epochs \
default 0.16700000000000000955 0.8499999999999997780    500
tanh     0.056000000000000000117 0.65000000000000002220    500
```

```
dropout 0.00100000000000000002 0.6999999999999995559 500
```

```
                train_loss                test_loss \
default 40.79299999999999926104 32.9089999999999891998
tanh    46.67099999999999937472 33.3459999999999653255
dropout 97.156000000000000591172 77.38200000000000500222
```

```
                train_loss_ep100            test_loss_ep100
default 49.661000000000000136424 35.26100000000000278533
tanh    47.72599999999999909051 33.1349999999999801048
dropout 428.0169999999999590727 394.2149999999997498890
```

1.10 Conclusiones

En este trabajo práctico pudimos ver como el gradiente descendente nos permite entrenar una red neuronal. Pudimos ver como la función de activación sigmoide es la mas adecuada para este tipo de problemas, y como el learning rate y la cantidad de epochs afectan la convergencia de la red.

También pudimos ver como la mascara de pesos no es una buena opción para este tipo de problemas, ya que no logra converger a una loss minima menor que la obtenida sin aplicar la mascara.

```
[ ]: nn_final = NeuralNet()
nn_final.fit(X_train, y_train, x_test=X_test, y_test=y_test, epochs=1000,
            lr=min_lr, eps=min_eps, verbose=True)

display(pd.Series(
    [
        y_test[0][0],
        nn_final.predict(X_test[0].reshape(1, -1))[0][0],
        nn_final.get_training_loss()[-1][0]
    ],
    index=["y_true", "y_pred", "train_loss"]
))

fig, ax = plt.subplots(figsize=(15, 10))
nn_final.plot_loss(ax=ax)
_ = ax.set_ylim(0, 100)
```

```
100%|      | 1000/1000 [00:04<00:00, 219.65it/s]
```

```
y_true      27.30000000000000071054
y_pred      27.23155726721838831850
train_loss  40.78857608365985498722
dtype: float64
```

